

## ▼ BANA274 Deep Learning Project

### Group 2 Project: Potato Chips Quality Control

**Team Members:** Mehak Katra, Aysouda Vatanparast, Kentrick Kepawitono, Urvi Vaidya, Moiz Nawab, Shahin Chinichian Moghaddam

```
#Setting up the model
from google.colab import drive
drive.mount('gdrive')

Mounted at gdrive

import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import scipy
from PIL import Image
from scipy import ndimage
import os


%matplotlib inline

from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, Flatten, Dropout, BatchNormalization
from tensorflow.keras import regularizers
from tensorflow.keras.callbacks import Callback, ModelCheckpoint, EarlyStopping
from tensorflow.keras.preprocessing import image
from tensorflow.keras.initializers import RandomNormal, he_normal, GlorotNormal
from sklearn.metrics import accuracy_score
from tensorflow.keras import optimizers

import tensorflow as tf
print(tf.__version__)

2.12.0

orig_im = Image.open('gdrive//My Drive//Colab Notebooks//Dataset//Train//Defective Chips//IMG_20210319_010422.jpg')
plt.imshow(orig_im)
plt.show()
```



```
#Assign training set with labels
defective_train_dir = "gdrive//My Drive//Colab Notebooks//Dataset//Augmented Training Set//Defected Chips"
non_defective_train_dir = "gdrive//My Drive//Colab Notebooks//Dataset//Augmented Training Set//Non-defect Chips"

X_train_orig = []
Y_train_orig = []
```

```

image_size = (64, 64)

# Load the defective training images and assign a label of 0
for filename in os.listdir(defective_train_dir):
    if filename.endswith(".jpg"):
        filepath = os.path.join(defective_train_dir, filename)
        image = Image.open(filepath)
        image = image.resize(image_size)
        image = np.array(image)
        X_train_orig.append(image)
        Y_train_orig.append(0)

# Load the non-defective training images and assign a label of 1
for filename in os.listdir(non_defective_train_dir):
    if filename.endswith(".jpg"):
        filepath = os.path.join(non_defective_train_dir, filename)
        image = Image.open(filepath)
        image = image.resize(image_size) #original size = (2500,2500)
        image = np.array(image)
        X_train_orig.append(image)
        Y_train_orig.append(1)

X_train_orig = np.array(X_train_orig)
Y_train_orig = np.array(Y_train_orig)

# Load the test data

# Define the paths to the directories with the test images
defective_test_dir = "gdrive//My Drive//Colab Notebooks//Dataset//Augmented Test Set//Defect Chips"
non_defective_test_dir = "gdrive//My Drive//Colab Notebooks//Dataset//Augmented Test Set//No defect Chips"

#Creating divisions for defective and non-defective images
filenames1 = os.listdir(defective_test_dir)
test_size1 = len(filenames1) // 2

filenames2 = os.listdir(non_defective_test_dir)
test_size2 = len(filenames2) // 2

test_filenames1 = filenames1[:test_size1]
crossval_filenames1 = filenames1[test_size1:]

test_filenames2 = filenames2[:test_size2]
crossval_filenames2 = filenames2[test_size2:]

X_test_orig = []
Y_test_orig = []
X_cv_orig = []
Y_cv_orig = []

image_size = (64, 64)

# Load the defective test & cross-validation images and assign a label of 0
for filename in test_filenames1:
    if filename.endswith(".jpg"):
        filepath = os.path.join(defective_test_dir, filename)
        image = Image.open(filepath)
        image = image.resize(image_size)
        image = np.array(image)
        X_test_orig.append(image)
        Y_test_orig.append(0)

for filename in crossval_filenames1:
    if filename.endswith(".jpg"):
        filepath = os.path.join(defective_test_dir, filename)
        image = Image.open(filepath)
        image = image.resize(image_size)
        image = np.array(image)
        X_cv_orig.append(image)
        Y_cv_orig.append(0)

# Load the non-defective test images and assign a label of 1
for filename in test_filenames2:
    if filename.endswith(".jpg"):
        filepath = os.path.join(non_defective_test_dir, filename)
        image = Image.open(filepath)
        image = image.resize(image_size)

```

```

image = np.array(image)
X_test_orig.append(image)
Y_test_orig.append(1)

for filename in crossval_filenames2:
    if filename.endswith(".jpg"):
        filepath = os.path.join(non_defective_test_dir, filename)
        image = Image.open(filepath)
        image = image.resize(image_size)
        image = np.array(image)
        X_cv_orig.append(image)
        Y_cv_orig.append(1)


X_test_orig = np.array(X_test_orig)
Y_test_orig = np.array(Y_test_orig)

X_cv_orig = np.array(X_cv_orig)
Y_cv_orig = np.array(Y_cv_orig)

index=40
plt.imshow(X_train_orig[index])
if Y_train_orig[index] == 1:
    print ("y = " + str(Y_train_orig[index]) + ", it's a non-defective potato chip.")
else:
    print ("y = " + str(Y_train_orig[index]) + ", it's a defective potato chip.")

```

y = 0, it's a defective potato chip.



```

m_train = X_train_orig.shape[0]
m_test = X_test_orig.shape[0]
m_cv = X_cv_orig.shape[0]
num_px = X_train_orig.shape[1]

print ("Number of training examples: m_train = " + str(m_train))
print ("Number of testing examples: m_test = " + str(m_test))
print ("Number of cross-validation examples: m_cv = " + str(m_cv))
print ("Height/Width of each image: num_px = " + str(num_px))
print ("Each image is of size: (" + str(num_px) + ", " + str(num_px) + ", 3)")
print ("train_set_x shape: " + str(X_train_orig.shape))
print ("train_set_y shape: " + str(Y_train_orig.shape))
print ("test_set_x shape: " + str(X_test_orig.shape))
print ("test_set_y shape: " + str(Y_test_orig.shape))
print ("cv_set_x shape: " + str(X_cv_orig.shape))
print ("cv_set_y shape: " + str(Y_cv_orig.shape))

```

```

Number of training examples: m_train = 8469
Number of testing examples: m_test = 384
Number of cross-validation examples: m_cv = 384
Height/Width of each image: num_px = 64
Each image is of size: (64, 64, 3)
train_set_x shape: (8469, 64, 64, 3)
train_set_y shape: (8469,)
test_set_x shape: (384, 64, 64, 3)
test_set_y shape: (384,)

```

```

cv_set_x shape: (384, 64, 64, 3)
cv_set_y shape: (384,)

# Reshape the training and test examples
train_set_x_flatten = X_train_orig.reshape(X_train_orig.shape[0],-1)
test_set_x_flatten = X_test_orig.reshape(X_test_orig.shape[0],-1)
cv_set_x_flatten = X_cv_orig.reshape(X_cv_orig.shape[0],-1)

print ("train_set_x_flatten shape: " + str(train_set_x_flatten.shape))
print ("train_set_y shape: " + str(Y_train_orig.shape))
print ("test_set_x_flatten shape: " + str(test_set_x_flatten.shape))
print ("test_set_y shape: " + str(Y_test_orig.shape))
print ("cv_set_x_flatten shape: " + str(cv_set_x_flatten.shape))
print ("cv_set_y shape: " + str(Y_cv_orig.shape))
print ("sanity check after reshaping: " + str(train_set_x_flatten[0:5,0]))

train_set_x_flatten shape: (8469, 12288)
train_set_y shape: (8469,)
test_set_x_flatten shape: (384, 12288)
test_set_y shape: (384,)
cv_set_x_flatten shape: (384, 12288)
cv_set_y shape: (384,)
sanity check after reshaping: [217 222 214 216 223]

print ("train_set_x shape: " + str((train_set_x_flatten.shape)))
print ("train_set_y shape: " + str(Y_train_orig.shape))
print ("test_set_x shape: " + str(test_set_x_flatten.shape))
print ("test_set_y shape: " + str(Y_test_orig.shape))
print ("cv_set_x shape: " + str(cv_set_x_flatten.shape))
print ("cv_set_y shape: " + str(Y_cv_orig.shape))

train_set_x shape: (8469, 12288)
train_set_y shape: (8469,)
test_set_x shape: (384, 12288)
test_set_y shape: (384,)
cv_set_x shape: (384, 12288)
cv_set_y shape: (384,)

```

#Standardizing dataset

```

train_set_x = train_set_x_flatten/255.
test_set_x = test_set_x_flatten/255.
cv_set_x = cv_set_x_flatten/255.

```

```

def get_test_accuracy(model, x_test, y_test):
    test_loss, test_acc = model.evaluate(x=x_test, y=y_test, verbose=0)
    print('accuracy: {acc:0.3f}'.format(acc=test_acc))

```

## Sequential Neural Network Model

```

model = Sequential([
    Dense(44, activation = 'leaky_relu', input_shape = (train_set_x.shape[1],)),
    Dense(25, activation = 'leaky_relu'),
    Dense(1, activation = 'sigmoid')
])

# compile the model
model.compile(optimizer=optimizers.legacy.Adam(learning_rate=0.00267),
              loss = tf.keras.losses.BinaryCrossentropy(from_logits=True),
              metrics = [tf.keras.metrics.BinaryAccuracy(threshold=0.5)])

```

## Saving best weights of NN model

```

early_stopping = EarlyStopping(monitor='val_binary_accuracy', patience=5, restore_best_weights=True)

#checkpoint

checkpoint_path = 'model_checkpoints/best_nn_weights'
checkpoint = ModelCheckpoint(filepath=checkpoint_path,
                             save_frequency='epoch',
                             save_weights_only=True,
                             monitor='val binary accuracy',

```

```

save_best_only=True,
verbose=1)

# fit the model
history = model.fit(train_set_x, Y_train_orig, epochs = 30, batch_size = 64, verbose = 2, callbacks=[checkpoint], validation_data=(cv_set_x, Y

133/133 - 2s - loss: 0.4792 - binary_accuracy: 0.7092 - val_loss: 0.4032 - val_binary_accuracy: 0.8021 - 2s/epoch - 13ms/step
Epoch 17/30

Epoch 17: val_binary_accuracy did not improve from 0.87500
133/133 - 2s - loss: 0.3956 - binary_accuracy: 0.8222 - val_loss: 0.4542 - val_binary_accuracy: 0.7839 - 2s/epoch - 14ms/step
Epoch 18/30

Epoch 18: val_binary_accuracy did not improve from 0.87500
133/133 - 1s - loss: 0.5050 - binary_accuracy: 0.8136 - val_loss: 2.4621 - val_binary_accuracy: 0.4792 - 1s/epoch - 11ms/step
Epoch 19/30

Epoch 19: val_binary_accuracy did not improve from 0.87500
133/133 - 1s - loss: 0.7294 - binary_accuracy: 0.7039 - val_loss: 0.5357 - val_binary_accuracy: 0.6823 - 1s/epoch - 10ms/step
Epoch 20/30

Epoch 20: val_binary_accuracy did not improve from 0.87500
133/133 - 1s - loss: 0.3931 - binary_accuracy: 0.8294 - val_loss: 0.4912 - val_binary_accuracy: 0.7578 - 1s/epoch - 10ms/step
Epoch 21/30

Epoch 21: val_binary_accuracy did not improve from 0.87500
133/133 - 1s - loss: 0.4448 - binary_accuracy: 0.7862 - val_loss: 0.5867 - val_binary_accuracy: 0.6667 - 1s/epoch - 11ms/step
Epoch 22/30

Epoch 22: val_binary_accuracy did not improve from 0.87500
133/133 - 2s - loss: 0.4302 - binary_accuracy: 0.7999 - val_loss: 0.3984 - val_binary_accuracy: 0.8255 - 2s/epoch - 17ms/step
Epoch 23/30

Epoch 23: val_binary_accuracy improved from 0.87500 to 0.89583, saving model to model_checkpoints/best_nn_weights
133/133 - 2s - loss: 0.3770 - binary_accuracy: 0.8314 - val_loss: 0.3484 - val_binary_accuracy: 0.8958 - 2s/epoch - 19ms/step
Epoch 24/30

Epoch 24: val_binary_accuracy did not improve from 0.89583
133/133 - 2s - loss: 0.3984 - binary_accuracy: 0.8098 - val_loss: 0.3301 - val_binary_accuracy: 0.8854 - 2s/epoch - 18ms/step
Epoch 25/30

Epoch 25: val_binary_accuracy did not improve from 0.89583
133/133 - 3s - loss: 0.3538 - binary_accuracy: 0.8545 - val_loss: 0.3559 - val_binary_accuracy: 0.8776 - 3s/epoch - 23ms/step
Epoch 26/30

Epoch 26: val_binary_accuracy did not improve from 0.89583
133/133 - 2s - loss: 0.3300 - binary_accuracy: 0.8623 - val_loss: 0.6367 - val_binary_accuracy: 0.7370 - 2s/epoch - 13ms/step
Epoch 27/30

Epoch 27: val_binary_accuracy did not improve from 0.89583
133/133 - 3s - loss: 0.4458 - binary_accuracy: 0.7949 - val_loss: 0.4267 - val_binary_accuracy: 0.7943 - 3s/epoch - 20ms/step
Epoch 28/30

Epoch 28: val_binary_accuracy did not improve from 0.89583
133/133 - 3s - loss: 1.8986 - binary_accuracy: 0.6506 - val_loss: 0.4202 - val_binary_accuracy: 0.8854 - 3s/epoch - 20ms/step
Epoch 29/30

Epoch 29: val_binary_accuracy did not improve from 0.89583
133/133 - 3s - loss: 0.4226 - binary_accuracy: 0.8242 - val_loss: 0.3683 - val_binary_accuracy: 0.8464 - 3s/epoch - 19ms/step
Epoch 30/30

Epoch 30: val_binary_accuracy did not improve from 0.89583
133/133 - 3s - loss: 0.4000 - binary_accuracy: 0.8235 - val_loss: 0.3471 - val_binary_accuracy: 0.8776 - 3s/epoch - 24ms/step

```

### Loss and metric(s) graph of NN model

```

# Plot the training and validation loss
frame = pd.DataFrame(history.history)
epochs = np.arange(len(frame))

fig = plt.figure(figsize=(12,4))

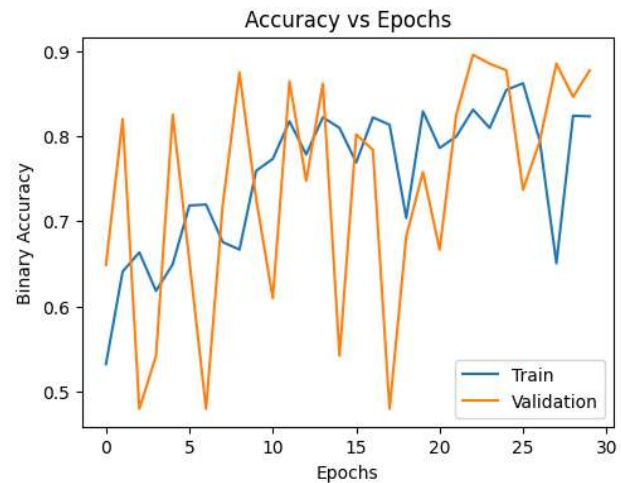
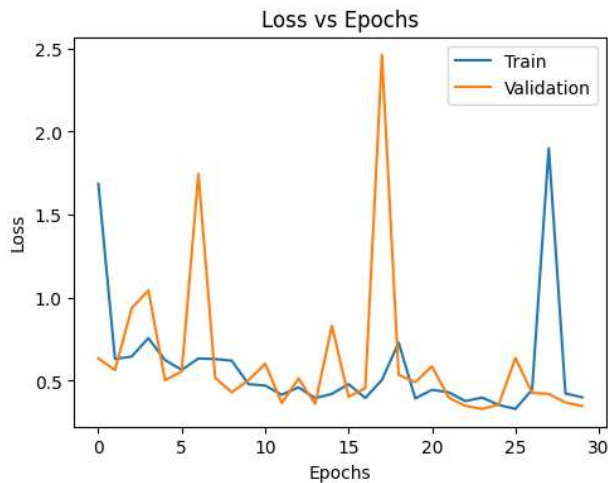
# Loss plot
ax = fig.add_subplot(121)
ax.plot(epochs, frame['loss'], label="Train")
ax.plot(epochs, frame['val_loss'], label="Validation")
ax.set_xlabel("Epochs")
ax.set_ylabel("Loss")

```

```
ax.set_title("Loss vs Epochs")
ax.legend()

# Accuracy plot
ax = fig.add_subplot(122)
ax.plot(epochs, frame['binary_accuracy'], label="Train")
ax.plot(epochs, frame['val_binary_accuracy'], label="Validation")
ax.set_xlabel("Epochs")
ax.set_ylabel("Binary Accuracy")
ax.set_title("Accuracy vs Epochs")
ax.legend()
```

<matplotlib.legend.Legend at 0x7f83d9e05810>



```
#test accuracy
get_test_accuracy(model, test_set_x, Y_test_orig)
```

```
/usr/local/lib/python3.10/dist-packages/keras/backend.py:5703: UserWarning: "`binary_crossentropy` received `from_logits=True`, but the
output, from_logits = _get_logits(
accuracy: 0.870
```

```
#cross-val accuracy
get_test_accuracy(model, cv_set_x, Y_cv_orig)
```

```
accuracy: 0.878
```

## CNN model

```
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, Flatten, BatchNormalization, Dropout, Conv2D, MaxPooling2D
from tensorflow.keras import regularizers

model1 = Sequential([
    Conv2D(filters=5, kernel_size=(3,3), strides=(1,1), padding='valid', activation = 'relu', input_shape=(64,64,3)),
    MaxPooling2D(pool_size=(2,2), strides=(1,1), padding='valid'),
    Conv2D(filters=5, kernel_size=(3,3), strides=(1,1), padding='valid', activation = 'relu'),
    MaxPooling2D(pool_size=(2,2), strides=(1,1), padding='valid'),
    Flatten(),
    Dense(120, activation = 'leaky_relu',kernel_regularizer=tf.keras.regularizers.l2(0.001)),
    BatchNormalization(),
    Dense(75, activation='leaky_relu',kernel_regularizer=tf.keras.regularizers.l2(0.001)),
    BatchNormalization(),
    Dense(1, activation='sigmoid', kernel_regularizer=tf.keras.regularizers.l2(0.001))
])
```

```
model1.summary()
```

```
Model: "sequential_1"
```

| Layer (type)    | Output Shape      | Param # |
|-----------------|-------------------|---------|
| conv2d (Conv2D) | (None, 62, 62, 5) | 140     |

```

max_pooling2d (MaxPooling2D (None, 61, 61, 5)      0
)
conv2d_1 (Conv2D)          (None, 59, 59, 5)      230
max_pooling2d_1 (MaxPooling2D) (None, 58, 58, 5)      0
flatten (Flatten)          (None, 16820)          0
dense_3 (Dense)             (None, 120)          2018520
batch_normalization (Batch Normalization) (None, 120)      480
dense_4 (Dense)             (None, 75)           9075
batch_normalization_1 (Batch Normalization) (None, 75)      300
dense_5 (Dense)             (None, 1)            76
=====
Total params: 2,028,821
Trainable params: 2,028,431
Non-trainable params: 390

```

---

```

model1.compile(optimizer=optimizers.legacy.Adam(learning_rate=0.00267),
               loss = tf.keras.losses.BinaryCrossentropy(from_logits=False),
               metrics = [tf.keras.metrics.BinaryAccuracy(threshold=0.5)])

```

### Saving the best weights for CNN model after training

```
early_stopping = EarlyStopping(monitor='val_binary_accuracy', patience=5, restore_best_weights=True)
```

```

checkpoint_path_cnn = 'model_checkpoints/checkpoint_CNN'
checkpoint = ModelCheckpoint(filepath=checkpoint_path_cnn,
                             save_frequency='epoch',
                             save_weights_only=True,
                             monitor='val_binary_accuracy',
                             save_best_only=True,
                             verbose=1)

```

```
cnn_history = model1.fit(X_train_orig, Y_train_orig, epochs = 15, batch_size = 64, verbose = 2, callbacks=[checkpoint], validation_data=(X_cv_
```

```

Epoch 1: val_binary_accuracy improved from -inf to 0.52344, saving model to model_checkpoints/checkpoint_CNN
133/133 - 46s - loss: 0.5806 - binary_accuracy: 0.9190 - val_loss: 5.5899 - val_binary_accuracy: 0.5234 - 46s/epoch - 342ms/step
Epoch 2/15

Epoch 2: val_binary_accuracy improved from 0.52344 to 0.73177, saving model to model_checkpoints/checkpoint_CNN
133/133 - 41s - loss: 0.2571 - binary_accuracy: 0.9685 - val_loss: 0.9651 - val_binary_accuracy: 0.7318 - 41s/epoch - 305ms/step
Epoch 3/15

Epoch 3: val_binary_accuracy improved from 0.73177 to 0.80469, saving model to model_checkpoints/checkpoint_CNN
133/133 - 53s - loss: 0.2224 - binary_accuracy: 0.9700 - val_loss: 0.4685 - val_binary_accuracy: 0.8047 - 53s/epoch - 396ms/step
Epoch 4/15

Epoch 4: val_binary_accuracy improved from 0.80469 to 0.98438, saving model to model_checkpoints/checkpoint_CNN
133/133 - 50s - loss: 0.1966 - binary_accuracy: 0.9760 - val_loss: 0.1645 - val_binary_accuracy: 0.9844 - 50s/epoch - 373ms/step
Epoch 5/15

Epoch 5: val_binary_accuracy did not improve from 0.98438
133/133 - 56s - loss: 0.1498 - binary_accuracy: 0.9790 - val_loss: 0.1886 - val_binary_accuracy: 0.9635 - 56s/epoch - 425ms/step
Epoch 6/15

Epoch 6: val_binary_accuracy improved from 0.98438 to 0.99219, saving model to model_checkpoints/checkpoint_CNN
133/133 - 47s - loss: 0.1459 - binary_accuracy: 0.9798 - val_loss: 0.0971 - val_binary_accuracy: 0.9922 - 47s/epoch - 351ms/step
Epoch 7/15

Epoch 7: val_binary_accuracy improved from 0.99219 to 0.99479, saving model to model_checkpoints/checkpoint_CNN
133/133 - 54s - loss: 0.1399 - binary_accuracy: 0.9835 - val_loss: 0.1031 - val_binary_accuracy: 0.9948 - 54s/epoch - 409ms/step
Epoch 8/15

Epoch 8: val_binary_accuracy did not improve from 0.99479

```

Epoch 9: val\_binary\_accuracy improved from 0.99479 to 1.00000, saving model to model\_checkpoints/checkpoint\_CNN  
 133/133 - 55s - loss: 0.1397 - binary\_accuracy: 0.9805 - val\_loss: 0.0916 - val\_binary\_accuracy: 1.0000 - 55s/epoch - 415ms/step  
 Epoch 10/15

Epoch 10: val\_binary\_accuracy did not improve from 1.00000  
 133/133 - 54s - loss: 0.1164 - binary\_accuracy: 0.9869 - val\_loss: 0.0888 - val\_binary\_accuracy: 0.9922 - 54s/epoch - 409ms/step  
 Epoch 11/15

Epoch 11: val\_binary\_accuracy did not improve from 1.00000  
 133/133 - 53s - loss: 0.1550 - binary\_accuracy: 0.9739 - val\_loss: 0.1043 - val\_binary\_accuracy: 0.9922 - 53s/epoch - 399ms/step  
 Epoch 12/15

Epoch 12: val\_binary\_accuracy did not improve from 1.00000  
 133/133 - 46s - loss: 0.1479 - binary\_accuracy: 0.9752 - val\_loss: 0.2643 - val\_binary\_accuracy: 0.9271 - 46s/epoch - 342ms/step  
 Epoch 13/15

Epoch 13: val\_binary\_accuracy did not improve from 1.00000  
 133/133 - 41s - loss: 0.1187 - binary\_accuracy: 0.9825 - val\_loss: 0.0783 - val\_binary\_accuracy: 0.9948 - 41s/epoch - 309ms/step  
 Epoch 14/15

Epoch 14: val\_binary\_accuracy did not improve from 1.00000  
 133/133 - 41s - loss: 0.1158 - binary\_accuracy: 0.9842 - val\_loss: 0.6392 - val\_binary\_accuracy: 0.7552 - 41s/epoch - 309ms/step  
 Epoch 15/15

Epoch 15: val\_binary\_accuracy did not improve from 1.00000  
 133/133 - 41s - loss: 0.1049 - binary\_accuracy: 0.9870 - val\_loss: 0.0873 - val\_binary\_accuracy: 0.9948 - 41s/epoch - 308ms/step

### Loss and metric(s) graphs

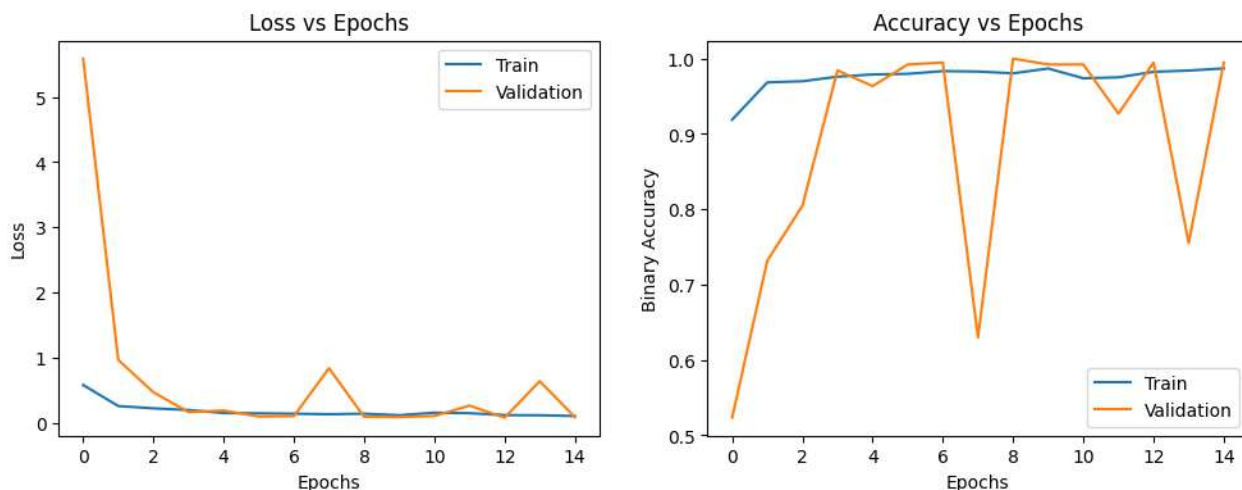
```
# Plot the training and validation loss
frame = pd.DataFrame(cnn_history.history)
epochs = np.arange(len(frame))

fig = plt.figure(figsize=(12,4))

# Loss plot
ax = fig.add_subplot(121)
ax.plot(epochs, frame['loss'], label="Train")
ax.plot(epochs, frame['val_loss'], label="Validation")
ax.set_xlabel("Epochs")
ax.set_ylabel("Loss")
ax.set_title("Loss vs Epochs")
ax.legend()

# Accuracy plot
ax = fig.add_subplot(122)
ax.plot(epochs, frame['binary_accuracy'], label="Train")
ax.plot(epochs, frame['val_binary_accuracy'], label="Validation")
ax.set_xlabel("Epochs")
ax.set_ylabel("Binary Accuracy")
ax.set_title("Accuracy vs Epochs")
ax.legend()
```

<matplotlib.legend.Legend at 0x7f841b8cc7f0>





```
#test accuracy
get_test_accuracy(model1, X_test_orig, Y_test_orig)

accuracy: 0.964

#test accuracy
get_test_accuracy(model1, X_cv_orig, Y_cv_orig)

accuracy: 0.995
```

### Loading weights of NN and CNN model

```
# NN model weights
model = Sequential([
    Dense(44, activation = 'leaky_relu', input_shape = (train_set_x.shape[1],)),
    Dense(25, activation = 'leaky_relu'),
    Dense(1, activation = 'sigmoid')
])

model.compile(optimizer=optimizers.legacy.Adam(learning_rate=0.00267),
              loss = tf.keras.losses.BinaryCrossentropy(from_logits=True),
              metrics = [tf.keras.metrics.BinaryAccuracy(threshold=0.5)])

model.load_weights(checkpoint_path)

<tensorflow.python.checkpoint.checkpoint.CheckpointLoadStatus at 0x7f841b594ee0>

# CNN model weights
model1 = Sequential([
    Conv2D(filters=5, kernel_size=(3,3), strides=(1,1), padding='valid', activation = 'relu', input_shape=(64,64,3)),
    MaxPooling2D(pool_size=(2,2), strides=(1,1), padding='valid'),
    Conv2D(filters=5, kernel_size=(3,3), strides=(1,1), padding='valid', activation = 'relu'),
    MaxPooling2D(pool_size=(2,2), strides=(1,1), padding='valid'),
    Flatten(),
    Dense(120, activation = 'relu',kernel_regularizer=tf.keras.regularizers.l2(0.001)),
    BatchNormalization(),
    Dense(75, activation='relu',kernel_regularizer=tf.keras.regularizers.l2(0.001)),
    BatchNormalization(),
    Dense(1, activation='sigmoid', kernel_regularizer=tf.keras.regularizers.l2(0.001))
])

model1.compile(optimizer=optimizers.legacy.Adam(learning_rate=0.00267),
              loss = tf.keras.losses.BinaryCrossentropy(from_logits=False),
              metrics = [tf.keras.metrics.BinaryAccuracy(threshold=0.5)])

model1.load_weights(checkpoint_path_cnn)

<tensorflow.python.checkpoint.checkpoint.CheckpointLoadStatus at 0x7f4cbaf48220>
```

### Random Test

```
# Select 5 images from test set

n = np.random.randint(0,X_test_orig.shape[0],5)
df=pd.Series()
for img in X_test_orig[n]:
    pic = Image.fromarray(img)
    df.loc[len(df)] = pic

<ipython-input-33-bd6137606c07>:4: FutureWarning: The default dtype for empty Series will be 'object' instead of 'float64' in a future v
df=pd.Series()

selected_images = test_set_x[n]
selected_images_cnn = X_test_orig[n]
selected_labels = Y_test_orig[n]

# Use previous model to predict labels
nn_predictions = model.predict(selected_images)
nn_predictions = np.where(nn_predictions > 0.5, 1, 0)
```

```
# Use CNN model to predict labels
cnn_predictions = model1.predict(selected_images_cnn)
cnn_predictions = np.where(cnn_predictions > 0.5, 1, 0)

# Display images with predicted labels
fig, axes = plt.subplots(nrows=1, ncols=5, figsize=(15, 3))

for i, ax in enumerate(axes):
    ax.imshow(selected_images_cnn[i])
    if nn_predictions[i] == 1:
        ndchip = "Non-Defective chip"
    else:
        ndchip = "Defective Chip"
    if cnn_predictions[i] == 1:
        cnnchip = "Non-Defective chip"
    else:
        cnnchip = "Defective Chip"
    ax.set_title(f'NN Pred: {ndchip} \n CNN Pred: {cnnchip}')
    ax.axis('off')

plt.tight_layout()
plt.show()
```

```
1/1 [=====] - 0s 28ms/step
1/1 [=====] - 0s 34ms/step
```

NN Pred: Defective Chip  
CNN Pred: Non-Defective chip




NN Pred: Defective Chip  
CNN Pred: Non-Defective chip



NN Pred: Defective Chip  
CNN Pred: Defective Chip



NN Pred: Non-Defective chip  
CNN Pred: Non-Defective chip



NN Pred: Defective Chip  
CNN Pred: Non-Defective chip



In our project, in the primitive model, we observed ~5% variance. We still tried to reduce this using regularization methods, but it increased the variance and affected the bias as well, thus introducing over-fitting and under-fitting under different scenarios.\_

So, we also checked the performance by developing a CNN model. It gave us the best results in comparison with the human performance. The variance and bias observed here are as expected- ~1-2%. This is acceptable according to industry standards. Even if a human tries to segregate the chips in real-time, ~5% human error can be expected when manufactured on a large scale.

Thus, to conclude, our CNN model gives the best performance. In the NN model, the base model performs the best without introducing weight initialization techniques/regularization techniques.